




(a) Optimisers vs LR, 2×64 (× = divergence)**(b) Architecture: deeper and wider help****(c) Learning curves by activation (4×64)****(d) Gradient size per layer at init (8 layers)**